

Dual-Branch Plant Disease Classification (Allen + Mufti)

This notebook builds a **dual-branch CNN** that fuses a custom convolutional pipeline with a pretrained ResNet34. It trains on two datasets:

- **Allen:** 3 classes (Healthy, Powdery, Rust), 1,530 images
- **Mufti:** 38 classes across 14 plants, ~70K train / ~17K test

Features:

1. Albumentations augmentations
2. Mixed-precision training (AMP)
3. Early stopping & LR scheduling
4. Logging & reproducibility
5. Confusion matrix & classification report
6. Gradio demo for live inference

1. Setup & Imports

Seed everything for reproducibility, set up logging, and import all required libraries.

```
pip install gradio
```

```
import os, random, logging
from pathlib import Path
```

```
import numpy as np
import pandas as pd
from PIL import Image
```

```
import torch import torch.nn as nn import
torch.nn.functional as F from torch.utils.data import
Dataset, DataLoader from torch.optim import AdamW from
torch.optim.lr_scheduler import ReduceLROnPlateau
```

```
import albumentations as A from
albumentations.pytorch import ToTensorV2
```

```
import torchvision from torchvision.models import * from
sklearn.metrics import confusion_matrix, classification_report
import matplotlib.pyplot as plt import seaborn as sns from
tqdm.auto import tqdm import gradio as gr
```

```
# Reproducibility def
seed_everything(seed=42):
    random.seed(seed)
    os.environ["PYTHONHASHSEED"] = str(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    torch.cuda.manual_seed_all(seed)
    torch.backends.cudnn.deterministic = True
    torch.backends.cudnn.benchmark = True
```

```
seed_everything(2025)
```

```
# Device device = torch.device("cuda" if torch.cuda.is_available()
else "cpu") print("Using device:", device)
```

```
# Logging logger = logging.getLogger("PlantDisease")
logger.setLevel(logging.INFO) ch = logging.StreamHandler()
ch.setFormatter(logging.Formatter("[%asctime)s] [%levelname)s] %(message)s"))
logger.addHandler(ch) fh = logging.FileHandler('plantdisease.log')
fh.setFormatter(logging.Formatter("[%asctime)s] [%levelname)s] %(message)s"))
logger.addHandler(fh)
```

```
import warnings
warnings.filterwarnings('ignore')
```

```
 Using device: cuda
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

2. Configuration

Define all paths and hyperparameters in a single Config class for easy adjustment.

```
class Config:
    # Dataset roots
    drive_root = Path("/content/drive/MyDrive")

    # Allen dataset has train/validation/test subfolders
    allen_train_dir = drive_root / "PDC/Allen/Train/Train"
    allen_val_dir = drive_root / "PDC/Allen/Validation/Validation"
    allen_test_dir = drive_root / "PDC/Allen/Test/Test"

    # Mufti dataset has an "Augmented" folder with train/valid, plus a separate test folder
    mufti_aug_dir = drive_root / "PDC/Mufti/New Plant Diseases Dataset(Augmented)"
    mufti_train_dir = mufti_aug_dir / "New Plant Diseases Dataset(Augmented)/train"
    mufti_val_dir = mufti_aug_dir / "New Plant Diseases Dataset(Augmented)/valid"
    mufti_test_dir = drive_root / "PDC/Mufti/test"

    # Image size & batch
    img_size = 64
    batch_size = 16

    # Training hyperparameters
    epochs = 20
    lr = 3e-4
    weight_decay = 1e-4
    patience = 3
    # for LR scheduler & early stopping
    cfg = Config()
```

```
logger.info(f"Allen dataset path: {cfg.allen_train_dir}")
logger.info(f"Allen train dir exists: {cfg.allen_train_dir.exists()}")
logger.info(f"Allen val dir exists: {cfg.allen_val_dir.exists()}")
logger.info(f"Allen test dir exists: {cfg.allen_test_dir.exists()}")
```

```
[2025-04-14 13:53:46,581][INFO] Allen dataset path: /content/drive/MyDrive/PDC/Allen/Train/Train
INFO:PlantDisease:Allen dataset path: /content/drive/MyDrive/PDC/Allen/Train/Train
[2025-04-14 13:53:49,203][INFO] Allen train dir exists: True
INFO:PlantDisease:Allen train dir exists: True
[2025-04-14 13:53:49,205][INFO] Allen val dir exists: True
INFO:PlantDisease:Allen val dir exists: True
[2025-04-14 13:53:49,209][INFO] Allen test dir exists: True
INFO:PlantDisease:Allen test dir exists: True
```

```
logger.info(f"Mufti dataset path: {cfg.mufti_train_dir}")
logger.info(f"Mufti train dir exists: {cfg.mufti_train_dir.exists()}")
logger.info(f"Mufti val dir exists: {cfg.mufti_val_dir.exists()}")
logger.info(f"Mufti test dir exists: {cfg.mufti_test_dir.exists()}")
```

```
[2025-04-14 13:53:49,214][INFO] Mufti dataset path: /content/drive/MyDrive/PDC/Mufti/New Plant Diseases Dataset(Augmented)/New Plant
INFO:PlantDisease:Mufti dataset path: /content/drive/MyDrive/PDC/Mufti/New Plant Diseases Dataset(Augmented)/New Plant Diseases Data
[2025-04-14 13:53:50,020][INFO] Mufti train dir exists: True
INFO:PlantDisease:Mufti train dir exists: True
[2025-04-14 13:53:50,022][INFO] Mufti val dir exists: True
INFO:PlantDisease:Mufti val dir exists: True
[2025-04-14 13:53:50,023][INFO] Mufti test dir exists: True
INFO:PlantDisease:Mufti test dir exists: True
```

3. Build Combined DataFrame

Scan both Allen and Mufti folders, prefix labels so they don't collide, and split into train/val/test.

```
def build_df(root: Path, prefix: str=""):
    """Return DataFrame with columns [filepath, label]."""
    rows = []
    if not root.exists():
        logger.warning(f"Path does not exist: {root}")
    return pd.DataFrame(rows)

    for cls in sorted(os.listdir(root)):
        cls_path = root / cls
        if not cls_path.is_dir():
            continue
        # Print the class path to check which folders it's iterating through
        print(f"Checking class path: {cls_path}")
        for img in cls_path.glob("*.jpg"):
            # Print the images found
            print(f"Found image: {img}")
            rows.append({"filepath": str(img), "label": prefix + cls})
```

```

# Print the number of rows added
print(f"Number of rows added: {len(rows)}")
return pd.DataFrame(rows)

# Allen dataset (3 classes) allen_train_df =
build_df(cfg.allen_train_dir, prefix="Allen_") allen_val_df =
build_df(cfg.allen_val_dir, prefix="Allen_") allen_test_df =
build_df(cfg.allen_test_dir, prefix="Allen_")

# Mufti dataset (38 classes) mufti_train_df =
build_df(cfg.mufti_train_dir, prefix="Mufti_") mufti_val_df =
build_df(cfg.mufti_val_dir, prefix="Mufti_") mufti_test_df =
build_df(cfg.mufti_test_dir, prefix="Mufti_")

# Combine datasets train_df = pd.concat([allen_train_df, mufti_train_df],
ignore_index=True).sample(frac=1, random_state=42) val_df = pd.concat([allen_val_df, mufti_val_df],
ignore_index=True).sample(frac=1, random_state=42) test_df = pd.concat([allen_test_df, mufti_test_df],
ignore_index=True).sample(frac=1, random_state=42)

# Sanity check: ensure DataFrames are not empty if
train_df.empty or val_df.empty or test_df.empty:
    logger.warning("One or more DataFrames are empty. Check your data paths and file formats.")
logger.warning(f"Allen train dir: {cfg.allen_train_dir} (exists={cfg.allen_train_dir.exists()})")
logger.warning(f"Allen val dir: {cfg.allen_val_dir} (exists={cfg.allen_val_dir.exists()})")
logger.warning(f"Allen test dir: {cfg.allen_test_dir} (exists={cfg.allen_test_dir.exists()})")
logger.warning(f"Mufti train dir: {cfg.mufti_train_dir} (exists={cfg.mufti_train_dir.exists()})")
logger.warning(f"Mufti val dir: {cfg.mufti_val_dir} (exists={cfg.mufti_val_dir.exists()})")
logger.warning(f"Mufti test dir: {cfg.mufti_test_dir} (exists={cfg.mufti_test_dir.exists()})")

# Encode labels: create mappings for consistency all_df =
pd.concat([train_df, val_df, test_df], ignore_index=True) labels =
sorted(all_df.label.unique()) label2idx = {l: i for i, l in
enumerate(labels)} idx2label = {i: l for l, i in
label2idx.items()}

# Map labels to indices for df_ in (train_df,
val_df, test_df): df_["label_idx"] =
df_.label.map(label2idx)

# Log final dataset sizes
logger.info(f"Final splits ► Train: {len(train_df)}, Val: {len(val_df)}, Test: {len(test_df)}")

```

 [Show hidden output](#)

```

logger.info(f"Train class distribution: \n{train_df['label'].value_counts()}")
logger.info(f"Val class distribution: \n{val_df['label'].value_counts()}")
logger.info(f"Test class distribution: \n{test_df['label'].value_counts()}")

```



Mufti Corn (maize) Common rust

477

```

Mufti_Corn_(maize)___Common_rust_ 477
Mufti_Grape___Black_rot 472
Mufti_Tomato___Leaf_Mold 470
Mufti_Corn_(maize)___healthy 465
Mufti_Tomato___Late_blight 463
Mufti_Peach___Bacterial_spot 459
Mufti_Tomato___Target_Spot 457
Mufti_Strawberry___healthy 456
Mufti_Cherry_(including_sour)___healthy 456
Mufti_Potato___healthy 456
Mufti_Blueberry___healthy 454
Mufti_Tomato___Tomato_mosaic_virus 448
Mufti_Raspberry___healthy 445
Mufti_Strawberry___Leaf_scorch 444
Mufti_Apple___Cedar_apple_rust 440
Mufti_Grape___Leaf_blight_(Isariopsis_Leaf_Spot) 440
Mufti_Tomato___Septoria_leaf_spot 436
Mufti_Tomato___Spider_mites_Two-spotted_spider_mite 435
Mufti_Squash___Powdery_mildew 434
Mufti_Grape___healthy 433
Mufti_Peach___healthy 432
Mufti_Tomato___Bacterial_spot 425
Mufti_Cherry_(including_sour)___Powdery_mildew 421
Mufti_Corn_(maize)___Cercospora_leaf_spot Gray_leaf_spot 410
Allen_Rust 20
Allen_Healthy 20
Allen_Powdery 20
Name: count, dtype: int64
[2025-04-14 13:55:22,327][INFO] Test class distribution:
label
Allen_Healthy 50
Allen_Rust 50
Allen_Powdery 50
Mufti_test 33
Name: count, dtype: int64
INFO:PlantDisease:Test class distribution:
label
Allen_Healthy 50
Allen_Rust 50
Allen_Powdery 50
Mufti_test 33
Name: count, dtype: int64

```

4. Dataset & DataLoader

Define a custom Dataset that applies Albumentations transforms. Then create PyTorch DataLoader s.

```

class PlantDataset(Dataset):
    def __init__(self, df, img_size, transforms=None):
        self.df = df.reset_index(drop=True)
        self.img_size = img_size
        self.transforms = transforms

    def __len__(self):
        return len(self.df)

    def __getitem__(self, idx):
        try:
            row = self.df.loc[idx]
            img_path = row.filepath
            # Log the image path to check for issues
            # print(f>Loading image: {img_path}")
            img = np.array(Image.open(img_path).convert("RGB"))
            if self.transforms:
                img = self.transforms(image=img)["image"]
            # Log the shape of the image after transformation
            # print(f"Image shape after transform: {img.shape}")
            return img, row.label_idx
        except Exception as e:
            # print(f"Error loading image {row.filepath}: {e}")
            # You could return a default image or raise the exception
            raise e

# Albumentations transforms
train_tfms = A.Compose([
    A.RandomResizedCrop(size=(cfg.img_size, cfg.img_size), scale=(0.8,1.0)), # Use size=(height, width)
    A.HorizontalFlip(0.5), A.VerticalFlip(0.5),
    A.ColorJitter(0.2,0.2,0.2,0.1),
    A.Normalize(), ToTensorV2(),
])
val_tfms = A.Compose([A.Resize(cfg.img_size, cfg.img_size), A.Normalize(), ToTensorV2()])

# Datasets
train_ds = PlantDataset(train_df, cfg.img_size, transforms=train_tfms)
val_ds = PlantDataset(val_df, cfg.img_size, transforms=val_tfms)

```

```
test_ds = PlantDataset(test_df, cfg.img_size, transforms=val_tfms)

# DataLoaders train_loader = DataLoader(train_ds, batch_size=cfg.batch_size, shuffle=True, num_workers=4,
pin_memory=True) val_loader = DataLoader(val_ds, batch_size=cfg.batch_size, shuffle=False, num_workers=4,
pin_memory=True) test_loader = DataLoader(test_ds, batch_size=cfg.batch_size, shuffle=False, num_workers=4,
pin_memory=True)
```

5. Dual-Branch Model Definition

Branch A: deep custom CNN

Branch B: frozen ResNet34 feature extractor Fusion:

concatenate features → classifier

```
import torch import torch.nn as nn
import torch.nn.functional as F
import torchvision.models as models

class DualBranchModel(nn.Module):
    def __init__(self, n_classes):
        super().__init__()
        # Branch A: custom CNN
        self.branch_a = nn.Sequential(
            nn.Conv2d(3,32,3,padding=1), nn.ReLU(), nn.Conv2d(32,32,3,padding=1), nn.ReLU(),
            nn.MaxPool2d(2), nn.Conv2d(32,64,3,padding=1), nn.ReLU(),
            nn.Conv2d(64,64,3,padding=1), nn.ReLU(), nn.MaxPool2d(2),
            nn.Conv2d(64,128,3,padding=1), nn.ReLU(), nn.Conv2d(128,128,3,padding=1), nn.ReLU(),
            nn.AdaptiveAvgPool2d(1), nn.Flatten(), nn.Linear(128,128), nn.ReLU()
        )
        # Branch B: pretrained ResNet34
        backbone = models.resnet34(pretrained=True)
        for p in backbone.parameters(): p.requires_grad=False
        self.branch_b = nn.Sequential(
            *list(backbone.children())[:-1], nn.Flatten(),
            nn.Linear(backbone.fc.in_features, 128),
            nn.ReLU()
        )
        # Fusion & classifier
        self.classifier = nn.Sequential(
            nn.Dropout(0.5),
            nn.Linear(256,256), nn.ReLU(), nn.Dropout(0.3),
            nn.Linear(256,n_classes)
        )

    def forward(self, x):
        a = self.branch_a(x)
        b = self.branch_b(x)
        fused = torch.cat([a,b], dim=1)
        return self.classifier(fused)

model = DualBranchModel(len(labels)).to(device)
logger.info(model)
```



(bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

```

    )
  )
  (5): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1,
affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (downsample): Sequential(
        (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (1): BasicBlock(
      (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1,
affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (2): BasicBlock(
      (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1,
affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (3): BasicBlock(
      (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1,
affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
)

```

6. Loss, Optimizer, Scheduler & AMP

Use label-smoothed cross-entropy, AdamW, ReduceLROnPlateau, and mixed precision.

```

criterion = nn.CrossEntropyLoss(label_smoothing=0.1) optimizer = AdamW(model.parameters(),
lr=cfg.lr, weight_decay=cfg.weight_decay) scheduler = ReduceLROnPlateau(optimizer, mode="min",
patience=cfg.patience, factor=0.5, verbose=True) scaler = torch.cuda.amp.GradScaler()

```



7. Training Loop with Early Stopping

Train for up to `cfg.epochs`, save the best model, and stop early if no improvement.

```
best_val_loss = float("inf") early_stop = 0 history = {}

{"train_loss": [], "train_acc": [], "val_loss": [], "val_acc": []}

for epoch in range(1, cfg.epochs+1):
    # - Training
    model.train()
    tloss, tcorrect = 0, 0
    for imgs,
    lbls in tqdm(train_loader, desc=f"Train {epoch}/{cfg.epochs}"):
        imgs, lbls = imgs.to(device),
        lbls.to(device)
        optimizer.zero_grad()
        with torch.cuda.amp.autocast():
            out =
            model(imgs)
            loss = criterion(out, lbls)
        scaler.scale(loss).backward()
        scaler.step(optimizer)
        scaler.update()

        preds = out.argmax(1)
        tloss
        += loss.item() * imgs.size(0)
        tcorrect += (preds==lbls).sum().item()

    # - Validation
    model.eval()
    vloss, vcorrect = 0, 0
    with
    torch.no_grad():
        for imgs, lbls in tqdm(val_loader, desc=f" Val
        {epoch}/{cfg.epochs}"):
            imgs, lbls = imgs.to(device),
            lbls.to(device)
            out = model(imgs)
            loss = criterion(out, lbls)
            preds =
            out.argmax(1)
            vloss += loss.item() *
            imgs.size(0)
            vcorrect += (preds==lbls).sum().item()

    # - Metrics
    train_loss =
    tloss/len(train_ds)
    train_acc =
    tcorrect/len(train_ds)
    val_loss =
    vloss/len(val_ds)
    val_acc =
    vcorrect/len(val_ds)

    history["train_loss"].append(train_loss)
    history["train_acc"].append(train_acc)
    history["val_loss"].append(val_loss)
    history["val_acc"].append(val_acc)

    logger.info(f"Epoch {epoch} ▶ Train: {train_loss:.4f}/{train_acc:.4f} | Val: {val_loss:.4f}/{val_acc:.4f}")
    scheduler.step(val_loss)

    # Early stopping
    if val_loss < best_val_loss:
        best_val_loss = val_loss
        torch.save(model.state_dict(), "best_model.pth")
    else:
        early_stop += 1
        if
        early_stop >= cfg.patience:
            logger.info("Early stopping triggered.")
            break
```

Plot Training Curves

```
plt.figure(figsize=(12,4)) plt.subplot(1,2,1)
plt.plot(history["train_loss"], '-o', label="Train Loss")
plt.plot(history["val_loss"], '-o', label="Val Loss")
plt.title("Loss"); plt.legend()

plt.subplot(1,2,2) plt.plot(history["train_acc"], '-o',
label="Train Acc") plt.plot(history["val_acc"], '-o',
label="Val Acc") plt.title("Accuracy"); plt.legend()
plt.show()
```

8. Test Set Evaluation

Load the best model, compute a confusion matrix and classification report.

```
# Load best weights
model.load_state_dict(torch.load("best_model.pth"))
model.eval()

y_true, y_pred = [], [] with torch.no_grad():    for
imgs, lbls in tqdm(test_loader, desc="Testing"):
    imgs = imgs.to(device)    out
    = model(imgs)    preds =
    out.argmax(1).cpu().numpy()
y_pred.extend(preds)
y_true.extend(lbls.numpy())
# Confusion matrix cm =
confusion_matrix(y_true, y_pred)
plt.figure(figsize=(10,8)) sns.heatmap(cm,
cmap="Blues", fmt="d") plt.title("Test
Confusion Matrix") plt.ylabel("True");
plt.xlabel("Predicted") plt.show()

# Classification report
print(classification_report(y_true, y_pred, target_names=labels))
```

10. Gradio Demo

Launch a simple web interface so judges can upload images and see top-3 predictions live.

```
def predict(img: np.ndarray):
    img_t = val_tfms(image=img)["image"].unsqueeze(0).to(device)
    with torch.no_grad():
        logits = model(img_t)    probs = F.softmax(logits,
dim=1)[0].cpu().numpy()    return {-idx2label[i]: float(probs[i]) for i
in range(len(labels))}

demo = gr.Interface(    fn=predict,    inputs=gr.Image(type="numpy", label="Upload-Leaf
Image"),    outputs=gr.Label(num_top_classes=3, label="Top-Predictions"),    title="Dual-
Branch Plant Disease Classifier",    description="Custom CNN + ResNet34 fusion trained on
Allen (3-class) & Mufti (38-class).",    examples=[["example1.jpg"], ["example2.jpg"]]
)
demo.launch()
```

 Running Gradio in a Colab notebook requires sharing enabled. Automatically setting `share=True` (you can turn this off by setting `s`


Colab notebook detected. To show errors in colab notebook, set debug=True in launch()


* Running on public URL: <https://1f02f169bcce485b23.gradio.live>

This share link expires in 72 hours. For free permanent hosting and GPU upgrades, run `gradio deploy` from the terminal in the worki

Dual-Branch Plant Disease Classifier

Custom CNN + ResNet34 fusion trained on Allen (3-class) & Mufti (38-class).




 Upload Leaf Image





Drop Image Here

- or -

Click to Upload






 Top Predictions



Flag

Clear
Submit

 Examples

